

# Modern Ways of Spatial Data Publication

Wouter Beek ([wouter@triply.cc](mailto:wouter@triply.cc)) and Laurens Rietveld ([laurens@triply.cc](mailto:laurens@triply.cc))

Triply (<http://triply.cc>)

v0.2

Almost every interesting dataset has some spatial component. Besides being prevalent, spatial relations – particularly geographical ones – tie the online to the offline world. As such, they provide a grounding of data stored in databases to the physical environment that is described in those databases.

Given the availability of Semantic Web services, Linked Datasets and Open Source web libraries we should be able to build a demonstration system that allows web programmers to build innovative applications on top of integrated Linked (Geo)datasets. Unfortunately, we found out that this is not (yet) the case.

## 1. Introduction

Earlier research by GeoNovum has resulted in a collection of lessons learned that describe in great detail the requirements of a modern spatial data publishing infrastructure. The purpose of the present report is to document our research findings based on this prior work in an attempt to answer the following research question:

How do the lessons learned meet the constraints (e.g., budgets) and capabilities (e.g., in-house know-how) of governmental organizations on the one hand, and of data users on the other?

While every governmental organization will be different, e.g., will be required to follow different rules and regulations depending on the domain or context in which it operates, it is possible to quantify the investment needed in order to build a Linked Geodata platform that implements the lessons learned.

## 2. First strategy: use a SotA triple store

When we started out building our geospatial Linked Data demonstrator we performed an industry-wide assessment of existing tools. For this we have consulted people from multiple domains. We first determined the leading query paradigm for geospatial Linked Data: this is GeoSPARQL [1]. We then determined the tool with the best GeoSPARQL support: this is Virtuoso<sup>1</sup>. We have loaded the data into an endpoint hosted at <http://sparql.geonovum.triply.cc/sparql>.

However, even Virtuoso does not support all geometries. Our data contains curves, resulting in the following error whenever a curve is encountered by the query engine:

```
Virtuoso 42000 Error GEO...: for after check of geo intersects, some
shape types (e.g., polygon rings and curves) are not yet supported
```

Here is an example of a query that gives the above warning (“Pairs of intersecting geometries that belong to the same resource.”):

```
PREFIX geo: <http://www.opengis.net/ont/geosparql#>
SELECT ?y1 ?y2
WHERE {
    ?x geo:asWKT ?y1 .
    ?x geo:asWKT ?y2
    FILTER (bif:st_intersects (?y1, ?y2, 0.1))
}
```

A second issue is that some queries do not terminate at all, as indicated by the following message:

```
Virtuoso S1T00 Error SR171: Transaction timed out
```

An example is an altered version of the above query (“Five pairs of dissimilar intersecting geometries that belong to the same resource.”):

```
PREFIX geo: <http://www.opengis.net/ont/geosparql#>
SELECT ?y1 ?y2 {
    ?x geo:asWKT ?y1 .
    ?x geo:asWKT ?y2
    FILTER (bif:st_intersects (?y1, ?y2, 0.1) && ?y1 != ?y2)
}
```

---

<sup>1</sup>See <https://github.com/openlink/virtuoso-opensource>

```
LIMIT 5
```

Thirdly, not all combinations of supported functions and supported shapes are supported. For instance, the following implements the geo query for our original use case:

```
PREFIX geo: <http://www.opengis.net/ont/geosparql#>
SELECT ?y (MIN(bif:st_distance(?y, bif:st_point(0, 52))) AS ?z)
WHERE {
    ?x1 geo:asWKT ?y
}
LIMIT 5
```

In this particular case, the problem is that distance can only be calculated between points but not between a point and a surface, as communicated by the following message:

```
Virtuoso 22023 Error GEO...: Function st_distance() expects a geometry
of type 1 as argument 0, not geometry of type 10242
```

The fourth and biggest problem is that queries that combine geo functions and graph relations take *very* long to compute. For testing purposes we have set the limit for query execution to 10,000 seconds, but some queries still do not succeed within that time-frame:

```
Virtuoso 42000 Error The estimated execution time 12774 (sec)
exceeds the limit of 10000 (sec).
```

We must note that both Virtuoso and StarDog are very good triple stores overall and that they are working on improving geodata and GeoSPARQL support. It is difficult to assess when those improvements will be good enough for the purpose of setting up a Linked Geodata service. It may also be the case that the out-of-the-box experience that we have tested can be improved through configuration changes that we are unaware of.

### 3. Second strategy: use a SotA Information Retrieval solution

We have corroborated our findings about the deficiency of existing triple stores with other developers. A common approach is to perform graph queries in a triple store and geospatial queries in a document store such as Solr. This approach results in a good performance and coverage of graph/SPARQL queries as well as good performance and coverage of geospatial queries. However, because the results come from different

backends, it is generally not possible to perform a GeoSPARQL query in which graph and geospatial components are intertwined.

## 4. Users have different needs & capacities (lesson 1)

One of the ways in which multiple groups of users can be addressed is by offering result set formats that they are familiar with. Our demonstrator exposes the following result set formats:

- GeoJSON
- JSON-LD 1.0
- N-Quads 1.1, N-Triples 1.1

N-Quads and N-Triples are, implementation-wise, the simplest RDF serialization formats available. They are supported by almost every RDF processor. Additional RDF serialization formats like Turtle 1.1, RDF/XML 1.1 and TRiG 1.1 are also widely available and are easy to add. We notice that the JavaScript RDF libraries<sup>2</sup> that are currently developed for client-side processing do not support, and will maybe never support, RDF/XML.

### 4.1. Header Dictionary Triples (HDT)

Another format that could be added is Header Dictionary Triples (HDT) [6]. These are currently used to power the Graph API to allow Basic Graph Pattern queries to be performed (similar to Linked Data Fragments (LDF) [9]). This format could be interesting for users who want to gather very large result sets.

Textual serialization formats work well for small results sets. For instance, if someone asks the ten nearest monuments then this can easily be returned in a text-based reply. However, some users have a large-scale use case, such as requesting *all* monuments in the Netherlands (e.g, with the purpose of data visualization) A textual result set of this size is difficult to process. With HDT the data is not only much smaller in size (as with regular compression) but can also be easily processed by the user.

---

<sup>2</sup>For more information, see the RDF JavaScript Libraries Group (<https://www.w3.org/community/rdfjs/>).

## 4.2. GeoJSON

Another format we expose is GeoJSON. The GeoJSON format is not yet fully standardized<sup>3</sup> and the current RFC [4] is not a standard in the traditional sense of the word, i.e., a definition of all and only GeoJSON constructs and their meaning. For instance, the current RFC does not give a formal grammar but relies on examples to convey GeoJSON syntax and semantics.

It is not very difficult to implement GeoJSON, which is a very flexible format that can easily be combined with other JSON formats. Specifically, we were interested in mixing GeoJSON into JSON-LD constructs. Some people have worked on this in 2015 under the name ‘geojson-ld’<sup>4</sup>, but development in that direction has stalled.

The biggest hurdle towards integrating GeoJSON and JSON-LD into one format is that JSON arrays are used in JSON-LD for abbreviated object term notation. However, GeoJSON uses JSON arrays to represent geometries (nested lists of floating point coordinates). This point will be addressed in JSON-LD 1.1<sup>5</sup>.

## 4.3. JSON-LD

JSON-LD 1.0 support was the most problematic result format to implement. It differs from Turtle-based serialization formats in that it does not translate a graph to a sequence of characters, but it instead performs transformations between RDF graphs and JSON trees. As such it has to marry requirements from both paradigms. In this sense JSON-LD is similar to RDF/XML which undertakes a graph to/from tree mapping as well.

JSON-LD is valid JSON, the primary data interchange format for web programmers. As such, JSON-LD is a good way of lowering the entry level for this user group. The JSON-LD format is relatively costly to generate and process, because not all aspects of RDF can be directly encoded in JSON. For this a *context* that steers the data transformation step needs to be defined (while generating) and applied (while processing). However, as long as JSON-LD is used for interchanging smaller chunks of data, the extra processing time is not an issue.

There are currently JSON-LD implementations for C#, Go, Java, JavaScript, PHP, Python, Ruby. Most notably support for C and C++, languages, in which many low-level database systems and libraries are written, is missing. Sadly, the top C++-based RDF processor, Raptor (<http://librdf.org/raptor/>), states on its web site that “JSON-LD

---

<sup>3</sup>GeoJSON is a proposed RFC standard as of August 2016.

<sup>4</sup>See <https://github.com/geojson/geojson-ld>

<sup>5</sup>See <https://github.com/json-ld/json-ld.org/issues/397>

is not supported - too complex to implement”.

Virtuoso does not support loading JSON-LD files<sup>6</sup> and ClioPatria does not support JSON-LD out-of-the-box either. We have written a partial implementation for generating JSON-LD and included it into library plRdf<sup>7</sup>.

#### 4.4. OGC standards (GML, WFS, WMS)

While GeoJSON addresses users from the web and geo domain, it may not address more advanced GIS users who may want to use more comprehensive formats standardized by the OGC. These formats include GML, WFS and WMS. The cost of integrating these OGC formats into a Linked Data platform are considerable, because they have their own vocabularies that need to be mapped to and from RDF. Luckily, the OGC and W3C are currently working on integrating their respective standards within the Spatial Data on the Web Working Group<sup>8</sup>. It is probably a good idea to wait until that Working Group has come up with a first (proposed) standard.

## 5. Findability (lesson 1a)

### 5.1. Vocabulary overview

One of the most difficult things for web programmers, when confronted with a Linked Geodata service, is to find out *what is in the data*. For non-Linked Data services this is usually not a problem: there is a limited number of entity types and relations that can be queried. For instance, a product review site may have users who write reviews for products. Products may have companies producing them and users may themselves have ratings to denote their trust level. That’s about it.

Linked Data is very different: the monument dataset that we started out with contains over a hundred unique relationships between entities belonging to dozens of different types. And this is only one dataset. It is inherently difficult to provide an overview of what is inside a large collection of Linked Data, in the same way in which it is impossible to provide an overview of what is in a large collection of the web. We have not found a satisfactory solution for this problem, which is widely recognized as a difficult problem in Linked Data. We do see several possible solutions but have not implemented them yet:

---

<sup>6</sup>See <https://github.com/openlink/virtuoso-opensource/issues/478>

<sup>7</sup>See [https://github.com/wouterbeek/plRdf/blob/master/prolog/jsonld/jsonld\\_build.pl](https://github.com/wouterbeek/plRdf/blob/master/prolog/jsonld/jsonld_build.pl)

<sup>8</sup>See [https://www.w3.org/2015/spatial/wiki/Main\\_Page](https://www.w3.org/2015/spatial/wiki/Main_Page)

- Faceted browsers map (part of) the schema to option lists. By selecting the options lists TODO
- Hierarchy visualization can be used to display the class and/or property hierarchy. StarDog uses this in their default data browser. Since our datasets currently do not contain hierarchies this options would first require a more detailed data model.

## 5.2. Data overview

- Various data visualization techniques have been proposed for Linked Data, but many of them only work with small data collections. E.g., spring embedding visualizations result in unwieldy visuals when they include more than a few thousand nodes.

## 5.3. Free text search

Datasets can be very large, so good search functions are required to let users find the needle in the haystack.

# 6. Keep it simple (lesson 1B)

Simplicity is notoriously difficult to achieve in Linked Data. The main reason for this is that simplicity is usually implemented by optimizing for a particular use case. Linked Data is about “the re-use of information in ways that are unforeseen by the publisher” [3]. However, this does not mean that existing approaches of Linked Geodata publishing cannot be simplified. We do observe that simplifications in Linked Geodata publishing are costly to implement. Simply changing something inside the User Interface is not enough: there are often conceptual reasons why certain things are difficult. We now give an example of a simplification we were able to implement.

## 6.1. Uniform & simple geodata representation

One of the ways in which geospatial data handling can be simplified is by enforcing a uniform representation format. The representation format must be generic enough to allow all representations that occur in the source datasets to be converted to it. At the same time, the uniform representation must allow data to be queried in a simple way. The following pattern is commonly used in our source datasets:

```

entity:x geosparql:defaultGeometry _:1 .
_:1 geosparql:asWKT "POLYGON(...)" ;
    rdf:type geosparql:Geometry .

```

The above three statements can be rewritten to a single statement conveying the same meaning:

```

entity:x def:geometry "POLYGON(...)"^^def:polygon .

```

This version has the following benefits:

- No introduction of an unnecessary blank node term (`_:1`).
- More explicit type information about the kind of geometry (`def:point` i.o. `geosparql:Geometry`).

These changes bring about the following usability improvements:

- One less level of nesting simplifies querying. For instance, we need one LDF request to retrieve the geometry of an entity instead of two.
- We can query for geometries of a specific type without having to parse geometry values. For instance, for one of our map views we want to show polygons but not points and lines. We achieve this with the following query:

```

?x def:geometry ?y FILTER (datatype(?y) = def:polygon)

```

We use Well-Known Text (WKT) as a uniform representation format, since it supports many different shapes (17) with a simple grammar<sup>9</sup> that is easy to implement. All other representations that appear in our source datasets can be converted to our single-triple representation. For instance, the following WGS84 representations

```

entity:y wgs84:lat "51.35"^^xsd:float ;
        wgs84:long "5.46"^^xsd:float .

```

is converted to

```

entity:y def:geometry "POINT(5.46 51.35)"^^def:polygon .

```

---

<sup>9</sup>See <http://svn.osgeo.org/postgis/trunk/doc/bnf-wkt.txt> for a Backus-Naur Form notation of the grammar.



## 7. Data transformation

Geodata, as it is currently published, does not contain explicit links. Links can be inferred from implicit cues by human domain experts, but not by automated and domain-independent means. This means that implementing lesson 2b (Foster to link everything with everything) requires a relatively large investment. We can distinguish the following issues with linking:

- Atomic terms are actually compound terms whose components and links are implicit. We solve this with *value unpacking* (Section 7.1).

The same value sometimes denotes the same thing. E.g., ‘Appingedam’ denotes a municipality. However, Appingedam is sometimes denoted by ‘GM0003’ (gemeentecode) and occurrences of the same string can denote different things (e.g., the area of Appingedam changes over time, the municipality of Appingedam is more than just the area of Appingedam, and a tourist uses ‘Appingedam’ to denote the center of Appingedam).

### 7.1. Value unpacking

A single value in the source dataset may describe multiple entities and relationships between them. For instance the following ‘buurtcode’ string BU00030002 denotes three entities and two spatial containment relations between them:

```
buurt:00030002  rdf:type      def:Buurt      .
gemeente:0003   geof:sfContains wijk:000300    ;
                rdf:type      def:Gemeente   .
wijk:000300     geof:sfContains buurt:00030002 ;
                rdf:type      def:Wijk       .
```

We call the conversion of a simple value to multiple entities and relations *value unpacking*. The idea behind the term is that domain experts have ‘packed’ meaning into encodings like BU00030002. In order to ‘open up’ this information to non-domain experts and machine processors, we have to ‘unpack’ the meaning again by using a grammar. For the above ‘buurtcode’ we have to construct the following grammar (written in Augmented Backus-Naur Form (ABNF) [5]):

```
buurt           := "BU" gemeente-code wijk-code buurt-code
buurt-code      := DIGIT DIGIT
gemeente        := "GM" gemeente-code
gemeente-code   := DIGIT DIGIT DIGIT DIGIT
```




Figure 1: Accuracy of natural language detection in RDF strings for the 10 most frequent languages. Strings of similar length are combined in buckets. The size of buckets increases logarithmically. Figure taken from [2].

```
wijk      := "WK" gemeente-code wijk-code(Wijk).
wijk-code := DIGIT DIGIT
```

When loading the data we have to tell the machine to *parse* values that appear in certain locations (either columns, tags or keys) by using the above grammar. The machine is able to automatically construct the encoded entities and relations for all conforming values. The high cost of value unpacking is caused by the fact that grammars are domain-dependent and by the fact that some grammars are more complex.

## 7.2. Value enrichment

Besides combining and splitting values, we sometimes find that a value is incomplete: it does not contain all the information we need in order to fully interpret it.

For instance, very many values are encoded in human languages. In order to do anything useful with these values we must at least know the language to which they belong. For this we run automatic language detection algorithms on the values. The accuracy for medium-length strings is known to be very high (Figure 1

The main benefit of running natural language detection on the current data collection is

to distinguish between Dutch and English strings.

### 7.3. Value combining

Sometimes multiple values can be combined into one aggregated value. This is specifically the case for datatypes. For example events are often described with a day, month and/or year column in the source day:

```
<EVENT-ID> | 1997 | Aug | 7
```

In order for dates to be comparable with one another in RDF they have to be converted to values of XML Schema Datatypes 1.1 [8]. Month names can be converted using a simple, one-to-one mapping. After that values can be converted to datatypes. The following

```
event:x def:day "07"^^xsd:gDay ;  
def:jaar "1997"^^xsd:gYear ;  
def:maand "8"^^xsd:gMonth .
```

would be combined into

```
event:x def:date "1997-08-07"^^xsd:date .
```

Because support for XML Schema Datatypes is very good overall, dates, times, durations, lengths, weights, etc. can all be compared with built-in functions. For instance the function `op:dateTime-less-than` is used by SPARQL 1.1 implementations to directly compare an event to all events that happened before it [7].

Value combining is only a good idea if the act of combining is lossless, i.e., no information is lost in the process. For instance, replacing the `foaf:givenName` and `foaf:familyName` properties with the combined `foaf:name` property is not an example of value combining, because it loses the cutoff point between the given and the family name.

### 7.4. Remove erroneous, empty and null values

## 8. Persistent IRIs (lesson 2c)

It is very difficult to ensure persistent IRIs. Most datasets have no IRIs at all, so we have to mint IRIs for each entity. This is easy when the source data has a key. If this is

not the case then

While enriching the data (Section 7) the number of entities increases. These requires IRIs as well.

We have found the Dutch Linked Data URI strategy to be very beneficial during this process. It distinguishes between IRIs for instance (ABOX) and schema (TBOX) entities.

## 9. Metadata

It is important for users to be able to get an overview of the datasets that are disseminated in a Linked Data service. While it is possible to insert a human-readable string description, this is not a sustainable solution.

It is a challenge to base

- `void:classes`
- `void:dataDump`
- `void:distinctObjects`
- `void:distinctSubjects`
- `void:feature`
- `void:properties`
- `void:triples`
- `void:vocabulary`

## 10. Quantifying the effort

There were also some ‘hidden costs’, i.e., either things that we expected to be very simple but that turned out to be very difficult, or things that we did not expect in the first place.

Firstly, we had to spend an enormous amount of time converting the various data formats (Section 7). After the conversion, we had to go through several iterations of transforming the data to improve its quality. While our transformed data has a much higher quality than the source data we started out with, we believe that there are still several more iterations needed to get the data to a quality level that allows generic (SPARQL) queries to be performed painlessly (i.e., without ad-hoc string manipulation).

Secondly, we expected to find a plethora of tools that allow Linked Data to be exposed to web programmers. After all, Linked Data is web data. We were very surprised to find

that this is not the case at all. In fact, there were many perfectly valid requests by the web programmers we worked with that could not be easily implemented by integrating some existing library.

### 10.1. Shortcuts

We believe that it is possible to speed up and reduce the cost of the process of developing a geospatial Linked Data web service by taking ‘shortcuts’.

**Points only** Many of the deficiencies of SotA triple stores can be worked around by reducing all geometries to 2D points. GeoSPARQL functions between points are mostly supported. The downside to this is that much of the geospatial information is lost and very many geospatial query (contains, intersects) cannot be performed at all. This also do not solve the problem that some queries never terminate or that many queries take too long.

**Start with Linked Data** In this research project we had to split development costs between implementing a web service and converting/transforming data. By starting out with Linked Data a big chunk of the costs can be saved (approx. 40% in our case).

## References

- [1] Robert Battle and Dave Kolas. GeoSPARQL: Enabling a geospatial Semantic Web. *Semantic Web Journal*, 3(4):355–370, 2011.
- [2] Wouter Beek, Filip Ilievski, Jeremy Debattista, Stefan Schlobach, and Jan Wielemaker. *Literally* better: Analyzing and improving the quality of literals. *Under submission*, 2016.
- [3] Tim Berners-lee, Yuhsin Chen, Lydia Chilton, Dan Connolly, Ruth Dhanaraj, James Hollenbach, Adam Lerer, and David Sheets. Tabulator: Exploring and analyzing linked data on the semantic web. In *Proceedings of the 3rd International Semantic Web User Interaction Workshop*, 2006.
- [4] H. Butler, M. Daly, A. Doyle, S. Gillies, S. Hagen, and T. Schaub. The GeoJSON format. *RFC 7946*, August 2016.
- [5] D. Crocker. Augmented BNF for syntax specifications: ABNF. *RFC 5234*, 2008.

Alias	IRI prefix
geo	<a href="http://www.opengis.net/ont/geosparql#">http://www.opengis.net/ont/geosparql#</a>
bif	<a href="http://www.openlinksw.com/schemas/bif#">http://www.openlinksw.com/schemas/bif#</a>
rdf	<a href="http://www.w3.org/1999/02/22-rdf-syntax-ns#">http://www.w3.org/1999/02/22-rdf-syntax-ns#</a>
rdfs	<a href="http://www.w3.org/2000/01/rdf-schema#">http://www.w3.org/2000/01/rdf-schema#</a>

Table 1: Aliases for commonly occurring IRI prefixes.

- [6] Javier D Fernández, Miguel A Martínez-Prieto, Claudio Gutiérrez, Axel Polleres, and Mario Arias. Binary RDF representation for publication and exchange (HDT). *Web Semantics: Science, Services and Agents on the World Wide Web*, 19:22–41, 2013.
- [7] Ashok Malhotra, Jim Melton, Norman Walsh, and Michael Kay. XQuery 1.0 and XPath 2.0 functions and operators (second edition). *W3C Recommendation*, April 2015.
- [8] David Peterson, Shudi (Sandy) Gao, Ashok Malhotra, C.M. Sperberg-McQueen, and Henry S. Thompson. XML Schema Definition Language (XSD) 1.1 part 2: Datatypes, April 2012.
- [9] Ruben Verborgh, Olaf Hartig, Ben De Meester, Gerald Haesendonck, Laurens De Vocht, Miel Vander Sande, Richard Cyganiak, Pieter Colpaert, Erik Mannens, and Rik Van de Walle. Querying datasets on the web with high availability. In *The Semantic Web–ISWC 2014*, pages 180–196. Springer, 2014.

## A. Used aliases